

From Zero to Cloud in 90 Days with Chef



Chef

Created by Brad Knowles
Based on materials from Chef Fundamentals OPS150-04.01 by Opscode, Inc.

- Name: Brad Knowles
- Company: ihiji, Inc.
- Experience:
 - System Administrator/Engineer/Consultant: 20+ years
 - Chef user: since 2011-09-14
- Contact
 - E-mail: bknowles@ihiji.com
 - Twitter: [@bradknowles](https://twitter.com/bradknowles)
 - github: [bknowles](https://github.com/bknowles)

What does ihiji do?

ihiji invasion is a web based solution that provides a secure gateway for residential electronics systems contractors & integrators to remotely monitor, service and maintain client's in-home electronic systems, including but not limited to, audio & video equipment, network devices, etc....

This comprises an on-premise appliance and back-end supporting cloud infrastructure with web portal for dealers to remotely monitor and manage equipment that is installed at the client site

What Can Be Monitored?

- Any IP Device
 - Ping
 - Well known ports and services (e.g., HTTP, FTP, Telnet, SSH, etc...)
 - Any device supporting SNMP
 - Control4 devices
- Any RS-232, Zigbee, AMX AxLink or Crestron Cresnet Device
- Custom monitoring modules can be written to support any IP based AV device

What Do The Appliances Look Like?



Problem

Company is growing fast, but old ad-hoc hand-built infrastructure can't scale nor can it be (easily) made fault-resilient

Company is growing fast, but old ad-hoc hand-built infrastructure can't scale nor can it be (easily) made fault-resilient



By Ferdinand Reus from Arnhem, Holland (Safari Uploaded by mangostar)
[CC-BY-SA-2.0], via [Wikimedia Commons](#)

Solution

- Rebuild from scratch
- In the cloud
- Must be scalable
- Must perform well
- Must be fault-resilient
- Must be very competitive re: cost/performance
- Using repeatable automated infrastructure management systems



Chef

Chef Can Help

Chef is designed to help manage this kind of complexity. You may have met already!

- Configuration management tool
- Systems integration framework
- API for infrastructure management

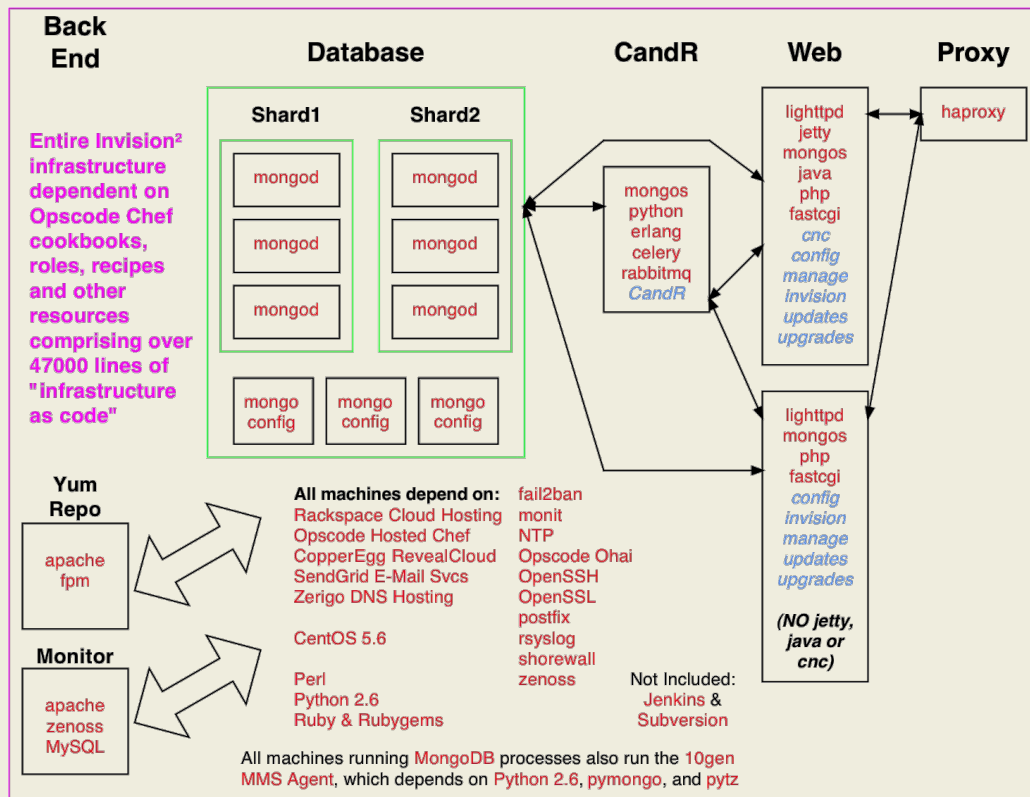
Problem

- I knew chef existed
 - But that was about it
- I had little prior exposure to CM tools
 - Only ever as a user of their services
 - Never as an engineer who implemented them
- I am **not** a developer

On the Plus Side

- I did have 20+ years of experience as a Unix System Engineer/Consultant
- I knew and trusted Matt Ray from his days at Zenoss
- Matt convinced us both that I could do this job with some training -- and some hands-on experience

Solution Diagram



Solution Logos



How Did We Get There?

- I Started work on August 16th, 2011
- In-depth interviews with management regarding current system design
- Reverse-engineer current systems
- Discover what needed to change and what could remain
- Working from first principles, deliver design documents:
 - High Level Architecture (40k foot)
 - Mid Level Services (10k foot)
 - Detail Level (one service @ 5k foot)
 - Project Goals
 - Rules of Thumb for Scalability
 - Next-generation Services and Service Names
- Use a couple of cloud instances to do some early testing

Training & Consulting

- Chef Fundamentals with Matt Ray
 - September 13th, 2011
- Rapid Prototype Infrastructure
 - Afternoon/Evening of September 13th, 2011
 - My first commit was on September 14th, 2011
 - Consulting complete by evening of September 15th, 2011

/trunk Developers

« [Development Statistics for /trunk](#)

Number of Developers:


5

| Author | Author Id | Changes | Lines of Code | Lines per Change |
|-----------------------------|-----------------------------|---------------|----------------|------------------|
| bknowles | bknowles | 2414 (82.6%) | 30311 (56.7%) | 12.5 |
| mray | mray | 311 (10.6%) | 20879 (39.1%) | 67.1 |
| srench | srench | 115 (3.9%) | 1376 (2.6%) | 11.9 |
| blantzsch | blantzsch | 80 (2.7%) | 869 (1.6%) | 10.8 |
| mmaniscalco | mmaniscalco | 1 (0.0%) | 8 (0.0%) | 8.0 |
| Totals | | 2921 (100.0%) | 53443 (100.0%) | 18.2 |


Last 12 Months

| Author | 9/2011 | 10/2011 | 11/2011 | 12/2011 | 1/2012 | 2/2012 | 3/2012 |
|-----------------------------|-------------------|------------------|------------------|------------------|-------------------|------------------|----------------|
| bknowles | 1838 (8.1%) | 1894 (77.0%) | 3114 (78.2%) | 7823 (92.0%) | 13509 (99.5%) | 2123 (97.3%) | 10 (100.0%) |
| mray | 20879 (91.8%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) |
| srench | 0 (0.0%) | 145 (5.9%) | 754 (18.9%) | 418 (4.9%) | 56 (0.4%) | 3 (0.1%) | 0 (0.0%) |
| blantzsch | 8 (0.0%) | 420 (17.1%) | 114 (2.9%) | 263 (3.1%) | 8 (0.1%) | 56 (2.6%) | 0 (0.0%) |
| mmaniscalco | 8 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) | 0 (0.0%) |
| Totals | 22733 (100.0%) | 2459 (100.0%) | 3982 (100.0%) | 8504 (100.0%) | 13573 (100.0%) | 2182 (100.0%) | 10 (100.0%) |


Commit Timeline



Contributed Lines of Code



Commit Heat Map



Patterns & Anti-Patterns

1. Make It Work, Then Make It Pretty
2. Publish Early, Publish Often
3. Use [Git](#)
4. Use [Github](#)
5. One Github Repo per Chef Cookbook
6. Use All Available Resources for Help
7. Use Code Reviews
8. Use Automated Testing
9. Use Continuous Integration/Deployment
10. Monitoring Is Just Continuous System/Service Testing
11. Logging is Monitoring Too
12. [ChaosMonkey](#) Is The (Current) Holy Grail

Make It Work, Then Make It Pretty

- Doing part one without part two is an **anti-pattern**
 - You will have jury-rigged a landmine in your infrastructure
 - Then foreverafter, you must always walk on eggshells around it
 - In twenty $\{time-units\}$ you won't remember what you did, why, or how and it **WILL** blow up in your face

Publish Early, Publish Often

- Helps encourage both parts of #1
 - Publishing something ugly is a negative feedback mechanism
- Encourages working in smaller increments that are more digestible

- All tools for Chef that assist with SCM workflow will assume git
 - Probably won't work with anything else
- As a Distributed VCS, git encourages frequent small commits, where tags and branches are trivially easy
- WhyGitIsBetterThanX

- Github is Facebook+LinkedIn for Developers & DevOps
 - It's how you show prospective employers what work you've done
 - It's how prospective employers show you what kind of work you might be doing
- Github is the largest Git community
 - Github works very hard to make it as easy as possible for people and groups to collaborate with each other
 - Git+Github is one of the most powerful tool combinations in your toolbox

One Github Repo per Chef Cookbook

- With Git and Github, collaboration occurs primarily at the repo level -- forking, sharing, remotes, committing, etc...
 - Each Chef cookbook should be largely standalone
 - Therefore, each Chef cookbook should be a separate repo
 - Can combine multiple repos together with [Librarian](#), [Braid](#), or git submodules

Use All Available Resources for Help

- No one can be an expert on everything
 - Whatever you're trying to do, someone else has probably already tried it
 - Don't stubbornly insist on re-inventing the wheel
- Use
 - Wiki: <http://wiki.opscode.com/>
 - Bug Tracker: <http://tickets.opscode.com/>
 - Community Cookbooks: <http://community.opscode.com/>
 - Community Mailing list: chef@lists.opscode.com
 - irc: #chef on irc.freenode.net
 - Hosted Chef:
 - Free for first five clients
 - Free access to Opscode Support: support@opscode.com

Use Code Reviews

- 80-90% of the time, I catch my code errors on proper inspection when I try to explain the code to someone else
- 80-90% of the remainder, my code errors are caught on inspection by the person who reviews my code
- If I skip code reviews, I lose the 95-99% error catch rate that having a programming partner would bring to the table

Use Automated Testing and Test Driven Development

- Chef [minitest](#) and [RSpec](#) are improving
 - Still beta quality code
 - Still not as well documented as they should be.
- [Cucumber-chef](#) is currently dependent on Amazon EC2
 - Needs more work (e.g., needs to support other cloud platforms)
 - Needs better documentation
 - [Test-Driven Infrastructure with Chef](#) is supposed to just be a "taste"
- Follow [Stephen Nelson-Smith \(@LordCope\)](#)
 - For all aspects of TDD/BDD, not just cucumber-chef

Use Continuous Integration/Deployment

- We use [Jenkins](#) but the tool you use is less important than doing CI in some fashion
 - [chef-jenkins](#) is the tool to use Jenkins to drive continuous deployment and synchronization of your Chef Environments from a git repository
- Chef is all about automating the management of the configuration of your infrastructure, and CI tools like Jenkins dovetail naturally with that

Monitoring Is Just Continuous System/Service Testing

- You need both internal and external monitoring
 - We use [Zenoss](#) and [monit](#) for internal monitoring
 - We use [CopperEgg](#) [RevealCloud](#) for external monitoring
- You need to monitor the systems, the applications, and the services
 - We use Zenoss and RevealCloud to monitor our systems
 - We use monit to monitor our applications
 - We use Zenoss to monitor our services

Logging is Monitoring Too

- Logging is just another way to monitor what is going on with your systems & applications
 - Forwarding & Gathering -- syslog-based (e.g., rsyslog & syslog-ng) and agent-based (incl. graylog agent, logstash, Splunk forwarders, etc...)
 - Analytics (e.g., graylog2, logstash, Splunk, etc...)
- Logging and log analysis is just as critical (if not more so) as monitoring
 - System, Service, and Application monitoring tell you **what** is happening
 - Log processing is more likely to be able to tell you **how** and **why**

ChaosMonkey Is The (Current) Holy Grail

- Failure **will** happen -- It's not a matter of **if**, but **when**
 - What if you could have some control over when failure occurs and how your systems respond?

Career Lessons

1. Know Yourself
2. Tools First
3. Value of Education, Dogfood, & #1 Revisited
4. DevOps vs. Politics
5. Family vs. Job, Value of Networking, & Rule of Three
6. More Value of Communities & Preparing to Speak/Teach
7. Location & Localization vs. Culturation & Distance
8. If You Don't Have the Proof, "It" Doesn't Exist
9. Choose Your Employers and Co-Workers Wisely

Know Yourself

- What is it that allows you to stay at a suboptimal job?
 - The People?
 - The Work?
 - The Money?
- For Me
 - People >> Work >> Money

Quote:

"If I had eight hours to chop down a tree, I'd spend six hours sharpening my axe. Give me six hours to chop down a tree and I will spend the first four sharpening the axe." -- Abraham Lincoln

Value of Education, Dogfood, & #1 Revisited

- Humans tend to not value something until it is gone or they are no longer there
- Eating your own dogfood is a sign that you are committed to succeed and you will bet your own business on your product
 - **However** -- Don't lose sight of who the real customer(s) is (are)

DevOps vs. Politics

- For DevOps to be successful, you **MUST** have complete buy-in from stem-to-stern
- As soon as a single person isn't 100% committed to the teamwork required, the whole corporate train starts going off the rails

Family vs. Job, Value of Networking, & Rule of Three

- Not Company vs. Home Life, but Company **AS** Family, vs. Company **AS** Job
 - If the company pitches itself as a family, but treats everyone as employees, you are on a path for disaster
- The Collective is Much More Resourceful and Much Smarter than the Individual
- However Long You Honestly Think it will Take, Multiply by Three

More Value of Communities & Preparing to Speak/Teach

- Experienced Smart People doing Intelligent Design with Open Source can lead to Novel Work
 - Which can lead to invitations to speak at conferences
 - Which can lead to meeting lots of other very experienced and smart people
- Preparing to speak at a conference or teach a tutorial tends to force me to really dig deep to completely and totally understand the material

Location & Localization vs. Culturation & Distance

- Significant physical distance makes work/life relationships hard
- Significant cultural and/or linguistic distance makes them harder
- Combine with economic downturn, and you may have the Perfect Storm

If You Don't Have the Proof, "It" Doesn't Exist

- You need to take a baseline when things are working well
 - If/when you need that baseline the time when you could have gotten it will have long since passed.

Choose Your Employers and Co-Workers Wisely

- Hiring = Talent Acquisition
 - Work for the kind of people you would want as friends regardless
 - Same for your co-workers
- If You Are Hiring
 - What kind of people you would want with you
 - If you were stuck in an elevator with them
 - ... for 24 hours
- Acquire the Right Talent
 - With the right personalities
 - With the right approaches to problem solving

Questions?

- E-mail: bknowles@ihiji.com
- Twitter: [@bradknowles](https://twitter.com/bradknowles)
- github: [bknowles](https://github.com/bknowles)